

Faster Discovery of Faster System Configurations with Spectral Learning

Vivek Nair, Tim Menzies, Xipeng Shen
NC State University, USA
vivekaxl, tim.menzies,
xipengshen@gmail.com

Norbert Siegmund, Sven Apel
University of Passau, Germany
norbert.siegmund, apel@uni-passau.de

ABSTRACT

Despite the huge spread and economical importance of configurable software systems, there is unsatisfactory support in utilizing the full potential of these systems with respect to finding performance-optimal configurations. Prior work on predicting the performance of software configurations suffered from either (a) requiring far too many sample configurations or (b) large variances in their predictions. Both these problems can be avoided using the **WHAT** spectral learner. **WHAT**'s innovation is the use of the spectrum (eigenvalues) of the distance matrix between the configurations of a configurable software system, to perform dimensionality reduction. Within that reduced configuration space, many closely associated configurations can be studied by executing only a few sample configurations. For the subject systems studied here, a few dozen samples yield accurate and stable predictors—less than 10% prediction error, with a standard deviation of less than 2%. When compared to the state of the art, our approach (a) requires 2 to 10 times fewer samples to achieve similar prediction accuracies, and (b) its predictions are more stable (i.e., have lower standard deviation). Furthermore, we demonstrate that predictive models generated by **WHAT** can be used by optimizers to discover system configurations that closely approach the optimal performance.

Categories/Subject Descriptors: D.2 [Software Engineering]; I.2 [Artificial Intelligence];

Keywords: Performance Prediction, Spectral Learning, Decision Trees, Search-Based Software Engineering, Sampling.

1. INTRODUCTION

Most software systems today are configurable. Despite the undeniable benefits of configurability, large configuration spaces challenge developers, maintainers, and users. In the face of hundreds of configuration options, it is difficult to keep track of the effects of individual configuration options and their mutual interactions. So, predicting the performance of individual system configurations or determining the optimal configuration is often more guess work than engineering. In their recent paper, Xu et al. documented the difficulties developers face with understanding the configuration

spaces of their systems [28]. As a result, developers tend to ignore over 5/6ths of the configuration options, which leaves considerable optimization potential untapped and induces major economic cost [28].

Addressing the challenge of performance prediction and optimization in the face of large configuration spaces, researchers have developed a number of approaches that rely on sampling and machine learning [7, 19, 24]. While gaining some ground, state-of-the-art approaches face two problems: (a) they require far too many sample configurations for learning or (b) they are prone to large variances in their predictions. For example, prior work on predicting performance scores using regression-trees had to compile and execute hundreds to thousands of specific system configurations [7]. A more balanced approach by Siegmund et al. is able to learn predictors for configurable systems [24] with low mean errors, but with large variances of prediction accuracy (e.g. in half of the results, the performance predictions for the Apache Web server were up to 50% wrong). Guo et al. [7] also proposed an incremental method to build a predictor model, which uses incremental random samples with steps equal to the number of configuration options (features) of the system. This approach also suffered from unstable predictions (e.g., predictions had a mean error of up to 22%, with a standard deviation of up to 46%). Finally, Sarkar et al. [19] proposed a projective-learning approach (using fewer measurements than Guo et al. and Siegmund et al.) to quickly compute the number of sample configurations for learning a stable predictor. However, as we will discuss, after making that prediction, the total number of samples required for learning the predictor is comparatively high (up to hundreds of samples).

The problems of large sample sets and large variances in prediction can be avoided using the **WHAT** spectral learner, which is our main contribution. **WHAT**'s innovation is the use of the spectrum (eigenvalues) of the distance matrix between the configurations of a configurable system, to perform dimensionality reduction. Within that reduced configuration space, many closely associated configurations can be studied by measuring only a few samples. In a number of experiments, we compared **WHAT** against the state-of-the-art approaches of Siegmund et al. [24], Guo et al. [7], and Sarkar et al. [19] by means of six real-world configurable systems: Berkeley DB, the Apache Web server, SQLite, the LLVM compiler, and the x264 video encoder. We found that **WHAT** performs as well or better than prior approaches, while requiring far fewer samples (just a few dozen). This is significant and most surprising, since some of the systems explored here have up to millions of possible configurations.

Overall, we make the following contributions:

- We present a novel sampling and learning approach for predicting the performance of software configurations in the face

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

Submitted to FSE'16 Seattle, WA, USA, 2016

© 2016 ACM. ISBN YYYY-YYYY-YY-YYY/YY/YYY...\$ZZ.00

DOI: XX.XXX/XXX+X

of large configuration spaces. The approach is based on a *spectral learner* that uses an approximation to the first principal component of the configuration space to recursively cluster it, relying only on a few points as representatives of each cluster.

- We demonstrate the practicality and generality of our approach by conducting experiments on six real-world configurable software systems (see Figure 1). The results show that our approach is more accurate (lower mean error) and more stable (lower standard deviation) than state-of-the-art approaches.
- We report on a comparative analysis of our approach and three state-of-the-art approaches, demonstrating that our approach outperforms previous approaches in terms of sample size and prediction stability. A key finding is the utility of the principal component of a configuration space to find informative samples from a large configuration space.

2. BACKGROUND & RELATED WORK

A configurable software system has a set X of Boolean configuration options,¹ also referred to as features or independent variables in our setting. We denote the number of features of system S as n . The configuration space of S can be represented by a Boolean space \mathbb{Z}_2^n , which is denoted by F . All valid configurations of S belong to a set V , which is represented by vectors \vec{C}_i (with $1 \leq i \leq |V|$) in \mathbb{Z}_2^n . Each element of a configuration represents a feature, which can either be *True* or *False*, based on whether the feature is selected or not. Each valid instance of a vector (i.e., a configuration) has a corresponding performance score associated to it.

The literature offers two approaches to performance prediction of software configurations: a *maximal sampling* and a *minimal sampling* approach: With *maximal sampling*, we compile all possible configurations and record the associated performance scores. Maximal sampling can be impractically slow. For example, the performance data used in this paper required 26 days of CPU time for measuring (and much longer, if we also count the time required for compiling the code prior to execution). Other researchers have commented that, in real world scenarios, the cost of acquiring the optimal configuration is overly expensive and time consuming [27].

If collecting performance scores of all configurations is impractical, *minimal sampling* can be used to intelligently select and execute just enough configurations (i.e., samples) to build a predictive model. For example, Zhang et al. [29] approximate the configuration space as a Fourier series, after which they can derive an expression showing how many configurations must be studied to build predictive models with a given error. While a theoretically satisfying result, that approach still needs thousands to hundreds of thousands of executions of sample configurations.

Another set of approaches are the four "additive" *minimal sampling* methods of Siegmund et al. [24]. Their first method, called feature-wise sampling (*FW*), is their basic method. To explain *FW*, we note that, from a configurable software system, it is theoretically possible to enumerate many or all of the valid configurations². Since each configuration (\vec{C}_i) is a vector of n Booleans, it is possible to use this information to isolate examples of how much each feature individually contributes to the total run time:

1. Find a pair of configurations \vec{C}_1 and \vec{C}_2 , where \vec{C}_2 uses exactly the same features as \vec{C}_1 , plus one extra feature f_i .
2. Set the run time $\Pi(f_i)$ for feature f_i to be the difference in the performance scores between \vec{C}_2 and \vec{C}_1 .
3. The run time for a new configuration \vec{C}_i (with $1 \leq i \leq |V|$) that has not been sampled before is then the sum of the run time of its features, as determined before:

$$\Pi(C_i) = \sum_{f_j \in C_i} \Pi(f_j) \quad (1)$$

When many pairs, such as \vec{C}_1, \vec{C}_2 , satisfy the criteria of point 1, Siegmund et al. used the pair that covers the *smallest* number of features. Their minimal sampling method, *FW*, compiles and executes only these smallest C_1 and C_2 configurations. Siegmund et al. also offers three extensions to the basic method, which are based on sampling not just the smallest \vec{C}_1, \vec{C}_2 pairs, but also any configurations with *interactions* between features. All the following minimal sampling policies compile and execute valid configurations selected via one of three heuristics:

PW (pair-wise): For each pair of features, try to find a configuration that contains the pair and has a minimal number of features selected.

HO (higher-order): Select extra configurations, in which three features, f_1, f_2, f_3 , are selected if two of the following pair-wise interactions exist: (f_1, f_2) and (f_2, f_3) and (f_1, f_3) .

HS (hot-spot): Select extra configurations that contain features that are frequently interacting with other features.

Guo et al. [7] proposed a progressive random sampling approach, which samples in steps of the number of features of the software system in question. They used the sampled configurations to train a regression tree, which is then used to predict the performance scores of other system configurations. The termination criterion of this approach is based on a heuristic, similar to the *PW* heuristics of Siegmund et al.

Sarkar et al. [19] proposed a cost model for predicting the effort (or cost) required to generate an accurate predictive model. The user can use this model to decide whether to go ahead and build the predictive model. This method randomly samples configurations and uses a heuristic based on feature frequencies as termination criterion. The samples are then used to train a regression tree; the accuracy of the model is measured by using a test set (where the size of the training set is equal to size of the test set). One of four projective functions (e.g., exponential) is selected based on how correlated they are to accuracy measures. The projective function is used to approximate the accuracy-measure curve, and the elbow point of the curve is then used as the optimal sample size. Once the optimal size is known, Sarkar et al. uses the approach of Guo et al. to build the actual prediction model.

The advantage of these previous approaches is that, unlike the results of Zhang et al., they require only dozens to hundreds of samples. Also, like our approach, they do not require to enumerate all configurations, which is important for highly configurable software systems. That said, as shown by our experiments (see Section 4), these approaches produce estimates with larger mean errors and partially larger variances than our approach. While sometimes the approach by Sarkar et al. results in models with (slightly) lower mean error rates, it still requires a considerably larger number of samples (up to hundreds, while **WHAT** requires only few dozen).

¹In this paper, we concentrate on Boolean options, as they make up the majority of all options; see Siegmund et al., for how to incorporate numeric options [23].

²Though, in practice, this can be very difficult. For example, in models like the Linux Kernel such an enumeration is practically impossible [20].

3. APPROACH

3.1 Spectral Learning

The minimal sampling method proposed in this paper is based on a spectral-learning algorithm that explores the spectrum (eigenvalues) of the distance matrix between configurations. In theory, such spectral learners are an appropriate method to handle noisy, redundant, and tightly inter-connected variables, for the following reasons: When data sets have many irrelevancies or closely associated data parameters d , then only a few eigenvectors e , $e \ll d$ are required to characterize the data. In that reduced space:

- Multiple inter-connected variables $i, j, k \subseteq d$ can be represented by a single eigenvector;
- Noisy variables from d are ignored, because they do not contribute to the signal in the data;
- Variables become (approximately) parallel lines in e space. For redundancies $i, j \in d$, we can ignore j since effects that change over j also change in the same way over i ;

That is, in theory, samples of configurations drawn via an eigenspace sampling method would not get confused by noisy, redundant, or tightly inter-connected variables. Accordingly, we expect predictions built from that sample to have lower mean errors and lower variances on that error.

Spectral methods have been used before for a variety of data mining applications [11]. Algorithms, such as PDDP [1], use spectral methods, such as principle component analysis (PCA), to recursively divide data into smaller regions. Software-analytics researchers use spectral methods (again, PCA) as a pre-processor prior to data mining to reduce noise in software-related data sets [26]. However, to the best of our knowledge, spectral methods have not been used before in software engineering as a basis of a minimal sampling method.

WHAT is somewhat different from other spectral learners explored in, for instance, image processing applications [21]. Work on image processing does not address defining a minimal sampling policy to predict performance scores. Also, a standard spectral method requires an $O(N^2)$ matrix multiplication to compute the components of PCA [10]. Worse, in the case of hierarchical division methods, such as PDDP, the polynomial-time inference must be repeated at every level of the hierarchy. Competitive results can be achieved using an $O(2N)$ analysis that we have developed previously [15], which is based on a heuristic proposed by Faloutsos and Lin [5] (which Platt has shown computes a Nyström approximation to the first component of PCA [17]).

Our approach inputs N (with $1 \leq |N| \leq |V|$) valid configurations $(\vec{C}), N_1, N_2, \dots$, and then:

1. Picks any point N_i ($1 \leq i \leq |N|$) at random;
2. Finds the point $West \in N$ that is furthest away from N_i ;
3. Finds the point $East \in N$ that is furthest from $West$.

The line joining $East$ and $West$ is our approximation for the first principal component. Using the distance calculation shown in Equation 2, we define δ to be the distance between $East$ and $West$. **WHAT** uses this distance (δ) to divide all the configurations as follows: The value x_i is the projection of N_i on the line running from $East$ to $West$ ³. We divide the examples based on the median

³The projection of N_i can be calculated in the following way:
 $a = dist(East, N_i); b = dist(West, N_i); x_i = \frac{a^2 - b^2 + \delta^2}{2\delta}$.

value of the projection of x_i . Now, we have two clusters of data divided based on the projection values (of N_i) on the line joining $East$ and $West$. This process is applied recursively on these clusters until a predefined stopping condition. In our study, the recursive splitting of the N_i 's stops when a sub-region contains less than $\sqrt{|N|}$ examples.

$$dist(x, y) = \begin{cases} \sqrt{\sum_i (x_i - y_i)^2} & \text{if } x_i \text{ and } y_i \text{ is numeric} \\ 0, & \text{if } x_i = y_i \\ 1, & \text{otherwise} \end{cases} \quad \text{if } x_i \text{ and } y_i \text{ is Boolean} \quad (2)$$

We explore this approach for three reasons:

- *It is very fast*: This process requires only $2|n|$ distance comparisons per level of recursion, which is far less than the $O(N^2)$ required by PCA [4] or other algorithms such as K-Means [8].
- *It is not domain-specific*: Unlike traditional PCA, our approach is general in that it does not assume that all the variables are numeric. As shown in Equation 2,⁴ we can approximate distances for both numeric and non-numeric data (e.g., Boolean).
- *It reduces the dimensionality problem*: This technique explores the underlying dimension (first principal component) without getting confused by noisy, related, and highly associated variables.

3.2 Spectral Sampling

When the above clustering method terminates, our sampling policy (which we will call S_1 :Random) is then applied:

Random sampling (S_1): compile and execute one configuration, picked at random, from each leaf cluster;

We use this sampling policy, because (as we will show later) it performs better than:

East-West sampling (S_2): compile and execute the *East* and *West* poles of the leaf clusters;

Exemplar sampling (S_3): compile and execute all items in all leaves and return the one with lowest performance score.

Note that S_3 is *not a minimal* sampling policy (since it executes all configurations). We use it here as one baseline against which we can compare the other, more minimal, sampling policies. In the results that follow, we also compare our sampling methods against another baseline using information gathered after executing all configurations.

3.3 Regression-Tree Learning

After collecting the data using one of the sampling policies (S_1 , S_2 , or S_3), as described in Section 3.2, we use a CART regression-tree learner [2] to build a performance predictor. Regression-tree learners seek the attribute-range split that most increases our ability to make accurate predictions. CART explores splits that divide N samples into two sets A and B , where each set has a standard deviation on the target variable of σ_1 and σ_2 . CART finds the “best” split defined as the split that minimizes $\frac{A}{N}\sigma_1 + \frac{B}{N}\sigma_2$. Using this best split, CART divides the data recursively.

In summary, **WHAT** combines:

⁴In our study, $dist$ accepts configurations (\vec{C}) and returns the distance between them. If x_i and $y_i \in \mathbb{R}^n$, then the distance function would be same as the standard Euclidean distance.

1. The FASTMAP method of Faloutsos and Lin [5];
2. A spectral-learning algorithm initially inspired by Boley’s PDDP system [1], which we modify by replacing PCA with FASTMAP (called “WHERE” in prior work [15]);
3. The sampling policy that explores the leaf clusters found by this recursive division;
4. The CART regression-tree learner that converts the data from the samples collected by sampling policy into a run-time prediction model [2].

That is,

$$\begin{aligned} \text{WHERE} &= \text{PDDP} - \text{PCA} + \text{FASTMAP} \\ \text{WHAT} &= \text{WHERE} + \text{SamplingPolicy} + \text{CART} \end{aligned}$$

This unique combination of methods has not been previously explored in the software-engineering literature.

4. EXPERIMENTS

All materials required for reproducing this work are available at <https://goo.gl/689Dve>.

4.1 Research Questions

We formulate our research questions in terms of the challenges of exploring large complex configuration spaces. Since our model explores the spectral space, our hypothesis is that only a small number of samples is required to explore the whole space. However, a prediction model built from a very small sample of the configuration space might be very inaccurate and unstable, that is, it may exhibit very large mean prediction errors and variances on the prediction error.

Also, if we learn models from small regions of the training data, it is possible that a learner will miss *trends* in the data between the sample points. Such trends are useful when building *optimizers* (i.e., systems that input one configuration and propose an alternate configuration that has, for instance, a better performance). Such optimizers might need to evaluate hundreds to millions of alternate configurations. To speed up that process, optimizers can use a *surrogate model*⁵ that mimics the outputs of a system of interest, while being computationally cheap(er) to evaluate [14]. For example, when optimizing performance scores, we might ask a CART for a performance prediction (rather than compile and execute the corresponding configuration). Note that such surrogate-based reasoning critically depends on how well the surrogate can guide optimization.

Therefore, to assess feasibility of our sampling policies, we must consider:

- Performance scores generated from our minimal sampling policy;
- The variance of the error rates when comparing predicted performance scores with actual ones;
- The optimization support offered by the performance predictor (i.e., can the model work in tandem with other off-the-shelf optimizers to generate useful solutions).

The above considerations lead to four research questions:

RQ1: *Can WHAT generate good predictions after executing only a small number of configurations?*

⁵Also known as response surface methods, meta models, or emulators.

Here, by “good” we mean that the predictions made by models that were trained using sampling with **WHAT** are as accurate, or more accurate, as those generated from models supplied with more samples.

RQ2: *Does less data used in building the models cause larger variances in the predicted values?*

RQ3: *Can “good” surrogate models (to be used in optimizers) be built from minimal samples?*

Note that **RQ2** and **RQ3** are of particular concern with our approach, since our goal is to sample as little as possible from the configuration space.

RQ4: *How good is WHAT compared to the state of the art of learning performance predictors from configurable software systems?*

To answer RQ4, we will compare **WHAT** against approaches presented by Siegmund et al. [24], Guo et al. [7], and Sarkar et al. [19].

System	LOC	Features	Configurations
Berkeley DB C Edition (BDBC)	219,811	18	2,560
Berkeley DB Java Edition (BDBJ)	42,596	32	400
Apache	230,277	9	192
SQLite	312,625	39	3,932,160
LLVM	47,549	11	1,024
x264	45,743	16	1,152

Figure 1: Subject systems used in the experiments.

4.2 Subject Systems

The configurable systems we used in our experiments are described in Figure 1. Note, with “predicting performance”, we mean predicting performance scores of the subject systems while executing test suites provided by the developers or the community, as described in Figure 1. To compare the predictions of our and prior approaches with actual performance measures, we use data sets that have been obtained by measuring *nearly all* configurations⁶. We say *nearly all* configurations, for the following reasoning: For all except one of our subject systems, the total number of valid configurations was tractable (192 to 2560). However, SQLite has

⁶<http://openscience.us/repo/performance-predict/cpm.html>

3,932,160 possible configurations, which is an impractically large number of configurations to test whether our predictions are accurate and stable. Hence, for SQLite, we use the 4500 samples for testing prediction accuracy and stability, which we could collect in one day of CPU time. Taking this into account, we will pay particular attention to the variance of the SQLite results.

4.3 Experimental Rig

RQ1 and **RQ2** require the construction and assessment of numerous runtime predictors from small samples of the data. The following rig implements that construction process.

For each configurable software system, we built a table of data, one row per valid configuration. We then ran all configurations of all software systems and recorded the performance scores (i.e., that are invoked by a benchmark). The exception is SQLite for which we measured only the configurations needed to detect interactions and additionally 100 random configurations to evaluate the accuracy of predictions. To this table, we added a column showing the performance score obtained from the actual measurements for each configuration.

Note that the following procedure ensures that we **never** test any prediction model on the data used to learn that model. Next, we repeated the following procedure 20 times (the figure of 20 repetitions was selected using the Central Limit Theorem): For each system in {BDBC, BDBJ, Apache, SQLite, LLVM, x264}

- Randomize the order of the rows in their table of data;
- For X in {10, 20, 30, ..., 90};
 - Let *Train* be the first X % of the data
 - Let *Test* be the rest of the data;
 - Pass *Train* to **WHAT** to select sample configurations;
 - Determine the performance scores associated with these configurations. This corresponds to a table lookup, but would entail compiling and executing a system configuration in a practical setting.
 - Using the *Train* data and their performance scores, build a performance predictor using CART.
 - Using the *Test* data, assess the accuracy of the predictor using the error measure of Equation 3 (see below).

The validity of the predictors built by the regression tree is verified on testing data. For each test item, we determine how long it *actually* takes to run the corresponding system configuration and compare the actual measured performance to the *prediction* from CART. The resulting prediction error is then computed using:

$$error = \frac{|predicted - actual|}{actual} * 100 \quad (3)$$

RQ2 requires testing the standard deviation of the prediction error rate. To support that test, we:

- Determine the X -th point in the above experiments, where all predictions stop improving (elbow point);
- Measure the standard deviation of the error at this point, across our 20 repeats.

As shown in Figure 2, all our results plateaued after studying $X = 40$ % of the valid configurations⁷. Hence to answer **RQ2**, we will compare all 20 predictions at $X = 40$ %.

⁷Just to clarify one frequently asked question about this work, we note that our rig “studies” 40 % of the data. We do not mean that our predictive models require accessing the performance scores from the 40 % of the data. Rather, by “study” we mean reflect on a sample of configurations to determine what minimal subset of that sample deserves to be compiled and executed.

RQ3 uses the learned regression tree as a *surrogate model* within an optimizer;

- Take $X = 40$ % of the configurations;
- Apply **WHAT** to build a CART model using some minimal sample taken from that 40 %;
- Use that CART model within some standard optimizer while searching for configurations with least runtime;
- Compare the faster configurations found in this manner with the fastest configuration known for that system.

This last item requires access to a ground truth of performance scores for a large number of configurations. For this experiment, we have access to that ground truth (since we have access to all system configurations, except for SQLite). Note that such a ground truth would not be needed when practitioners choose to use **WHAT** in their own work (it is only for our empirical investigation).

For the sake of completeness, we explored a range of optimizers seen in the literature in this second experiment: DE [25], NSGA-II [3], and our own GALE [12, 30] system. Normally, it would be reasonable to ask why we used those three, and not the hundreds of other optimizers described in the literature [6, 9]. However, as shown below, all these optimizers in this domain exhibited very similar behavior (all found configurations close to the best case performance). Hence, the specific choice of optimizer is not a critical variable in our analysis.

5. RESULTS

5.1 RQ1

Can WHAT generate good predictions after executing only a small number of configurations?

Figure 2 shows the mean errors of the predictors learned after taking X % of the configurations, then asking **WHAT** and some sampling method (S_1 , S_2 , and S_3) to (a) find what configurations to measure; then (b) asking CART to build a predictor using these measurements. The horizontal axis of the plots shows what X % of the configurations are studied; the vertical axis shows the mean relative error (from Equation 3). In that figure:

- The \times — \times lines in Figure 2 show a *baseline* result where data from the performance scores of 100 % of configurations were used by CART to build a runtime predictor.
- The other lines show the results using the sampling methods defined in Section 3.2. Note that these sampling methods used runtime data only from a subset of 100 % of the performance scores seen in configurations from 0 to X %.

In Figure 2, *lower* y-axis values are *better* since this means lower prediction errors. Overall, we find that:

- Some software systems exhibit large variances in their error rate, below $X = 40$ % (e.g., BDBC and BDBJ).
- Above $X = 40$ %, there is little effect on the overall change of the sampling methods.
- Mostly, S_3 shows the highest overall error, so that it cannot be recommended.
- Always, the \times — \times baseline shows the lowest errors, which is to be expected since predictors built on the baseline have access to all data.




Figure 2: Errors of the predictions made by WHAT with four different sampling policies. Note that, on the y-axis, lower errors are better.

- We see a trend that the error of S_1 and S_2 are within 5% of the baseline results. Hence, we can recommend these two minimal sampling methods.

Figure 3 provides information about which of S_1 or S_2 we should recommend. This figure displays data taken from the $X = 40\%$ point of Figure 2 and displays how many performance scores of configurations are needed by our sub-sampling methods (while reflecting on the configurations seen in the range $0 \leq X \leq 40$). Note that:

- S_3 needs up to thousands of performance-score points, so it cannot be recommended as minimal-sampling policy;
- S_2 needs twice as much performance-score information as S_1 (S_2 uses two samples per leaf cluster while S_1 uses only one).
- S_1 needs performance-score information on only a few dozen (or less) configurations to generate the predictions with the lower errors seen in Figure 2.

Combining the results of Figure 2 and Figure 3, we conclude that:




Figure 3: Comparing evaluations of different sampling policies. We see that the number of configurations evaluated for S_2 is twice as high as S_1 , as it selects 2 points from each cluster, where as S_1 selects only 1 point.

S_1 is our preferred spectral sampling method. Furthermore, the answer to **RQ1** is “yes”, because applying **WHAT**, we can (a) generate runtime predictors using just a few dozens of sample performance scores; and (b) these predictions have error rates within 5% of the error rates seen if predictors are built from information about all performance scores.

5.2 RQ2

Do less data used in building the models cause larger variances in the predicted values?

Two competing effects can cause increased or decreased variances in runtime predictions. The less we sample the configuration space, the less we constrain model generation in that space. Hence, one effect that can be expected is that models learned from too few samples exhibit large variances. But, a compensating effect can be introduced by sampling from the spectral space since that space contains fewer confusing or correlated variables than the raw configuration space.

Figure 4 reports which one of these two competing effects are dominant. Figure 2 shows that after some initial fluctuations, after seeing $X = 40\%$ of the configurations, the variances in prediction errors reduces to nearly zero.

Hence, we answer **RQ2** with “no”: Selecting a small number of samples does not necessarily increase variance (at least to say, not in this domain).

5.3 RQ3

Can “good” surrogate models (to be used in optimizers) be built from minimal samples?

The results of answering **RQ1** and **RQ2** suggest to use **WHAT** (with S_1) to build runtime predictors from a small sample of data. **RQ3** asks if that predictor can be used by an optimizer to infer what other configurations correspond to system configurations with fast performance scores. To answer this question, we ran a random set




Figure 4: Standard deviations seen at various points of Figure 2.

of 100 configurations, 20 times, and related that baseline to three optimizers (GALE [12], DE [25] and NSGA-II [3]) using their default parameters.

When these three optimizers mutated existing configurations to suggest new ones, these mutations were checked for validity. Any mutants that violated the system’s constraints (e.g., a feature excluding another feature) were rejected and the survivors were “evaluated” by asking the CART surrogate model. These evaluations either rejected the mutant or used it in generation $i + 1$, as the basis for a search for more, possibly better mutants.

Figure 5 shows the configurations found by three optimizers projected onto the ground truth of the performance scores of nearly all configurations (see Section 4.2). Again note that, while we use that ground truth for the validation of these results, our optimizers used only a small part of that ground-truth data in their search for the fastest configurations (see the **WHAT** + S_1 results of Figure 3).

The important feature of Figure 5 is that all the optimized configurations fall within 1% of the fastest configuration according to the ground truth (see all the left-hand-side dots on each plot). Table 1 compares the performance of the optimizers used in this study. Note that the performances are nearly identical, which leads to the following conclusions:




Figure 5: Solutions found by GALE, NSGA-II and DE (shown as points) laid against the ground truth (all known configuration performance scores). It can be observed that all the optimizers can find the configuration with lower performance scores.

The answer to **RQ3** is “yes”: For optimizing performance scores, we can use surrogates built from few runtime samples. The choice of the optimizer does not critically effect this conclusion.

5.4 RQ4

How good is WHAT compared to the state of the art of learning performance predictors from configurable software systems?

We compare **WHAT** with the three state-of-the-art predictors proposed in the literature [24], [7], [19], as discussed in Section 2. Note that all approaches use regression-trees as predictors, except Siegmund’s approach, which uses a regression function derived using linear programming.

The bars of Figure 6 show the mean error rate, the standard deviation of the error rate, and the mean percentage of total config-

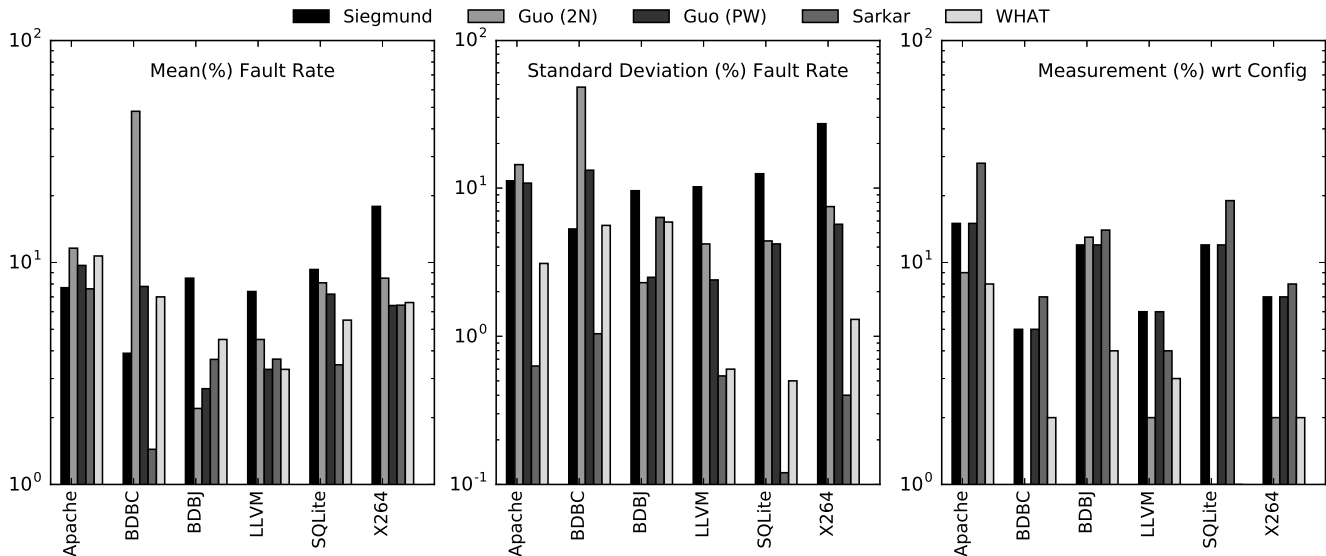


Figure 6: Comparison between WHAT and the state-of-the-art approaches regarding mean error, standard deviation, and the percentage of configurations used for training the model.

Table 1: The table shows how the minimum performance scores as found by the learners GALE, NSGA-II, and DE, vary over 20 repeated runs. Mean values are denoted μ and IQR denotes the 25th–75th percentile. A low IQR suggests that the surrogate model build by WHAT is stable and can be utilized by off the shelf optimizers to find performance-optimal configurations.

Dataset	Searcher					
	GALE		DE		NSGAI	
	Mean	IQR	Mean	IQR	Mean	IQR
Apache	870	0	840	0	840	0
BDBC	0.363	0.004	0.359	0.002	0.354	0.005
BDBJ	3139	70	3139	70	3139	70
LLVM	202	3.98	200	0	200	0
SQLite	13.1	0.241	13.1	0	13.1	0.406
X264	248	3.3	244	0.003	244	0.05

urations used in 30 repeats of the different approaches. Note that the y-axis of that figure is a logarithmic scale so, within each plot:

- Differences near the bottom are very small differences;
- Differences near the top are very large differences;

As seen in the left and middle plots of Figure 6, the *FW* approach of Siegmund et al. (i.e., the sampling approach using the fewest number of configurations) often has the highest error rate and the highest standard deviation on that error rate. Hence, we cannot recommend this method or, if one wishes to use this method, we recommend using the other sampling heuristics (e.g., HO, HS) to make more accurate predictions (but at the cost of much more measurements). Moreover, the size of the standard deviation of this method causes further difficulties in estimating which configurations are those exhibiting a large prediction error.

As to the approach of Guo et al. (with PW), this does not stand out on any of our measurements. Its error results are within 1 % of

WHAT ; its standard deviation are usually larger; and it requires much more data than WHAT .

In terms of the number of measured samples required to build a model, the right-hand-side plot of Figure 6 shows that WHAT requires the fewest samples except for two cases: the approach of Guo et al. (with 2N) working on BDBC and LLVM. In both these cases, the mean error and standard deviation on the error estimate is larger than WHAT (see the \square bars in the left and middle plots of Figure 6). Furthermore, in the case of BDBC, the error values are $\mu = 14\%$, $\sigma = 13\%$, which are much larger than WHAT’s error rates of $\mu = 6\%$, $\sigma = 5\%$.

Although the approach of Sarkar et al. produces an error rate that is sometimes less than the one of WHAT (see the left-hand-side of Figure 6), it requires the most number of measurements. Moreover, WHAT’s accuracy is close to Sarkar’s approach (1 % to 2 % difference). Hence, we cannot recommend this approach, too.

The right-hand-side of Figure 6 shows the *percent* of required measurements. Table 2 shows the same expressed as absolute values. We see that most state-of-the-art approaches often require many more samples than WHAT . Using those fewest numbers of samples, WHAT has within 1 to 2 % of the lowest standard deviation rates and within 1 to 2 % of lowest error rates. The exception is Sarkar’s approach, which has 5 % lower mean error rates (in BDBC, see the left-hand-side plot of Figure 6). However, as shown in right-hand-side of Table 2, Sarkar’s approach needs nearly three times more measurements than WHAT (191 vs 64 samples). Given the overall reduction of the error is small (5 % difference between Sarkar and WHAT in mean error), the overall cost of tripling the data-collection cost is often not feasible in a practical context and might not justify the small additional benefit in accuracy.

Hence, we answer RQ4 with “yes”, since WHAT yields predictions that are similar to or more accurate than prior work, while requiring fewer samples.

Table 2: Comparison of the number of the samples required with the state of the art. The grey colored cells indicate the approach which has the lowest number of samples. We notice that WHAT and Guo (2N) uses less data compared to other approaches. The high fault rate of Guo (2N) accompanied with high variability in the predictions makes WHAT our preferred method.

Dataset	Samples				
	Siegmund	Guo (2N)	Guo (PW)	Sarkar	WHAT
Apache	29	181	29	55	16
BDBC	139	36	139	191	64
BDBJ	48	52	48	57	16
LLVM	62	22	64	43	32
SQLite	566	78	566	925	64
X264	81	32	81	93	32

6. RELIABILITY AND VALIDITY

Reliability refers to the consistency of the results obtained from the research. For example, how well independent researchers could reproduce the study? To increase external reliability, this paper has taken care to either clearly define our algorithms or use implementations from the public domain (SciKitLearn) [16]. Also, all the data used in this work is available on-line in the PROMISE code repository and all our algorithms are on-line at github.com/ai-se/where.

Validity refers to the extent to which a piece of research actually investigates what the researcher purports to investigate [22]. *Internal validity* checks if the differences found in the treatments can be ascribed to the treatments under study.

One internal validity issue with our experiments is the choice of *training and testing* data sets discussed in Figure 1. Recall that while all our learners used the same *testing* data set, our untuned learners were only given access to *training* data.

Another internal validity issues is *instrumentation*. The very low μ and σ error values reported in this study are so small that it is reasonable to ask whether they are due to some instrumentation quirk, rather than due to using a clever sample strategy:

- Our low μ values are consistent with prior work (e.g. [19]);
- As to our low σ values, we note that, when the error values are so close to 0%, the standard deviation of the error is “squeezed” between zero and those errors. Hence, we would expect that experimental rigs that generate error values on the order of 5% and Equation 3 should have σ values of $0 \leq \sigma \leq 5$ (e.g., like those seen in our introduction).

Regarding SQLite, we cannot measure all possible configurations in reasonable time. Hence, we sampled only 100 configurations to compare prediction and actual performance values. We are aware that this evaluation leaves room for outliers. Also, we are aware that measurement bias can cause false interpretations [15]. Since we aim at predicting performance for a special workload, we do not have to vary benchmarks.

We aimed at increasing the *external validity* by choosing software systems from different domains with different configuration mechanisms and implemented with different programming languages. Furthermore, the systems used are deployed and used in the real world. Nevertheless, assuming the evaluations to be automatically transferable to all configurable software systems is not fair. To further strengthen external validity, we run the model (generated by WHAT + S_1) against other optimizers, such as NSGA-II and dif-

ferential evolution [25]. That is, we validated whether the learned models are not only applicable for GALE style of perturbation. In Table 1, we see that the models developed are valid for all optimizers, as all optimizers are able to find the near optimal solutions.

7. RELATED WORK

In 2000, Shi and Maik [21] claimed the term “spectral clustering” as a reference to their normalized cuts image segmentation algorithm that partitions data through a spectral (eigenvalue) analysis of the Laplacian representation of the similarity graph between instances in the data.

In 2003, Kamvar et al. [11] generalized that definition saying that “spectral learners” were any data-mining algorithm that first replaced the raw dimensions with those inferred from the spectrum (eigenvalues) of the affinity (a.k.a. distance) matrix of the data, optionally adjusted via some normalization technique).

Our clustering based on first principal component splits the data on a approximation to an eigenvector, found at each recursive level of the data (as described in §3.1). Hence, this method is a “spectral clusterer” in the general Kamvar sense. Note that, for our data, we have not found that Kamvar’s normalization matrices are needed.

Regarding sampling, there are a wide range of methods know as experimental designs or designs of experiments [18]. They usually rely on fractional factorial designs as in the combinatorial testing community [13].

Furthermore, there is a recent approach that learns *performance-influence models* for configurable software systems [23]. While this approach can handle even numeric features, it has similar sampling techniques for the Boolean features as reported in their earlier work [24]. Since we already compared to that earlier work and do not consider also numeric features, we did not compare our work to performance-influence models.

8. CONCLUSIONS

Configurable software systems today are widely used in practice, but expose challenges regarding finding performance-optimal configurations. State-of-the-art approaches require too many measurements or are prone to large variances in their performance predictions. To avoid these shortcomings, we have proposed a fast spectral learner, called **WHAT**, along with three new sampling techniques. The key idea of **WHAT** is to explore the configuration space with eigenvalues of the features used in a configuration to determine exactly those configurations for measurement that reveal key performance characteristics. This way, we can study many closely associated configurations with only a few measurements.

We evaluated our approach on six real-world configurable software systems borrowed from the literature. Our approach achieves similar to lower error rates, while being stable when compared to the state of the art. In particular, with the exception of Berkeley DB, our approach is more accurate than the state-of-the-art approaches by Siegmund et al. [24] and Guo et al. [7]. Furthermore, we achieve a similar prediction accuracy and stability as the approach by Sarkar et al [19], while requiring a far smaller number of configurations to be measured. We also demonstrated that our approach can be used to build cheap and stable surrogate prediction models, which can be used by off-the-shelf optimizers to find the performance-optimal configuration.

9. REFERENCES

- [1] Daniel Boley. Principal direction divisive partitioning. *Data mining and knowledge discovery*, 2(4):325–344, 1998.
- [2] Leo Breiman, Jerome Friedman, Charles J Stone, and Richard A Olshen. *Classification and regression trees*. CRC press, 1984.
- [3] Kalyanmoy Deb, Amrit Pratap, Sameer Agarwal, and T Meyarivan. A fast and elitist multiobjective genetic algorithm: NSGA-II. *IEEE Transactions on Evolutionary Computation*, 6(2):182–197, 2002.
- [4] Qian Du and James E Fowler. Low-complexity principal component analysis for hyperspectral image compression. *International Journal of High Performance Computing Applications*, 22(4):438–448, 2008.
- [5] Christos Faloutsos and King-Ip Lin. Fastmap: A fast algorithm for indexing, data-mining and visualization of traditional and multimedia datasets. volume 24. ACM, 1995.
- [6] Roger Fletcher. *Practical methods of optimization*. John Wiley & Sons, 2013.
- [7] Jianmei Guo, Krzysztof Czarnecki, Sven Apel, Norbert Siegmund, and Andrzej Wasowski. Variability-aware performance prediction: A statistical learning approach. In *IEEE/ACM 28th International Conference on Automated Software Engineering*, pages 301–311. IEEE, 2013.
- [8] Greg Hamerly. Making k-means even faster. Society for Industrial and Applied Mathematics.
- [9] Mark Harman, S Afshin Mansouri, and Yuanyuan Zhang. Search-based software engineering: Trends, techniques and applications. *ACM Computing Surveys*, 45(1):11, 2012.
- [10] Alexander Ilin and Tapani Raiko. Practical approaches to principal component analysis in the presence of missing values. *The Journal of Machine Learning Research*, 11:1957–2000, 2010.
- [11] Kamvar Kamvar, Sepandar Sepandar, Klein Klein, Dan Dan, Manning Manning, and Christopher Christopher. Spectral learning. In *International Joint Conference of Artificial Intelligence*. Stanford InfoLab, 2003.
- [12] Joseph Krall, Tim Menzies, and Misty Davies. Gale: Geometric active learning for search-based software engineering. *IEEE Transactions on Software Engineering*, 41(10):1001–1018, 2015.
- [13] D Richard Kuhn, Raghu N Kacker, and Yu Lei. *Introduction to combinatorial testing*. CRC press, 2013.
- [14] Ilya Gennadyevich Loshchilov. *Surrogate-assisted evolutionary algorithms*. PhD thesis, 2013.
- [15] Tim Menzies, Andrew Butcher, David Cok, Andrian Marcus, Lucas Layman, Forrest Shull, Burak Turhan, and Thomas Zimmermann. Local versus global lessons for defect prediction and effort estimation. *IEEE Transactions on Software Engineering*, 39(6):822–834, 2013.
- [16] Fabian Pedregosa, Gaël Varoquaux, Alexandre Gramfort, Vincent Michel, Bertrand Thirion, Olivier Grisel, Mathieu Blondel, Peter Prettenhofer, Ron Weiss, Vincent Dubourg, et al. Scikit-learn: Machine learning in python. *The Journal of Machine Learning Research*, 12:2825–2830, 2011.
- [17] John Platt. Fastmap, Metricmap, and Landmark MDS are all nystrom algorithms. pages 261–268. Society for Artificial Intelligence and Statistics, 2005.
- [18] Friedrich Pukelsheim. *Optimal Design of Experiments*, volume 50. Society for Industrial and Applied Mathematics, 1993.
- [19] Atri Sarkar, Jianmei Guo, Norbert Siegmund, Sven Apel, and Krzysztof Czarnecki. Cost-efficient sampling for performance prediction of configurable systems. In *30th IEEE/ACM International Conference on Automated Software Engineering*, pages 342–352. IEEE, 2015.
- [20] Abdel Salam Sayyad, Joe Ingram, Tim Menzies, and Hany Ammar. Scalable product line configuration: A straw to break the camel’s back. In *IEEE/ACM 28th International Conference on Automated Software Engineering*, pages 465–474. IEEE, 2013.
- [21] Jianbo Shi and Jitendra Malik. Normalized cuts and image segmentation. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 22(8):888–905, 2000.
- [22] Janet Siegmund, Norbert Siegmund, and Sven Apel. Views on internal and external validity in empirical software engineering. In *Proceedings of the 37th International Conference on Software Engineering*, pages 9–19. IEEE, 2015.
- [23] Norbert Siegmund, Alexander Grebhahn, Sven Apel, and Christian Kästner. Performance-influence models for highly configurable systems. In *Proceedings of the 2015 10th Joint Meeting on Foundations of Software Engineering*, pages 284–294. ACM, 2015.
- [24] Norbert Siegmund, Sergiy S Kolesnikov, Christian Kästner, Sven Apel, Don Batory, Marko Rosenmüller, and Gunter Saake. Predicting performance via automated feature-interaction detection. In *Proceedings of the 34th International Conference on Software Engineering*, pages 167–177. IEEE Press, 2012.
- [25] Rainer Storn and Kenneth Price. Differential evolution—a simple and efficient heuristic for global optimization over continuous spaces. *Journal of global optimization*, 11(4):341–359, 1997.
- [26] Christopher Theisen, Kim Herzig, Patrick Morrison, Brendan Murphy, and Laurie Williams. Approximating attack surfaces with stack traces. In *Proceedings of the 37th International Conference on Software Engineering*, pages 199–208. IEEE Press, 2015.
- [27] Gary M Weiss and Ye Tian. Maximizing classifier utility when there are data acquisition and modeling costs. *Data Mining and Knowledge Discovery*, 17(2):253–282, 2008.
- [28] Tianyin Xu, Long Jin, Xuepeng Fan, Yuanyuan Zhou, Shankar Pasupathy, and Rukma Talwadder. Hey, you have given me too many knobs!: Understanding and dealing with over-designed configuration in system software. In *Proceedings of the 2015 10th Joint Meeting on Foundations of Software Engineering*, pages 307–319. ACM, 2015.
- [29] Yi Zhang, Jianmei Guo, Eric Blais, and Krzysztof Czarnecki. Performance prediction of configurable software systems by fourier learning. In *30th IEEE/ACM International Conference on Automated Software Engineering*, pages 365–373. IEEE, 2015.
- [30] Marcela Zuluaga, Guillaume Sergent, Andreas Krause, and Markus Püschel. Active learning for multi-objective optimization. In *Proceedings of the 30th International Conference on Machine Learning*, pages 462–470, 2013.